



4.4 Lecture Summary

4 Dataflow Synchronization and Pipelining

4.4 Pipeline Parallelism

Lecture Summary: In this lecture, we studied how point-to-point synchronization can be used to build a one-dimensional *pipeline* with p tasks (stages), T_0, \dots, T_{p-1} . For example, three important stages in a *medical imaging* pipeline are *denoising*, *registration*, and *segmentation*.

We performed a simplified analysis of the *WORK* and *SPAN* for pipeline parallelism as follows. Let n be the number of input items and p the number of stages in the pipeline, $WORK = n \times p$ is the total work that must be done for all data items, and $CPL = n + p - 1$ is the *span* or *critical path length* for the pipeline. Thus, the ideal parallelism is $PAR = WORK / CPL = np / (n + p - 1)$. This formula can be validated by considering a few boundary cases. When $p = 1$, the ideal parallelism degenerates to $PAR = 1$, which confirms that the computation is sequential when only one stage is available. Likewise, when $n = 1$, the ideal parallelism again degenerates to $PAR = 1$, which confirms that the computation is sequential when only one data item is available. When n is much larger than p ($n \gg p$), then the ideal parallelism approaches $PAR = p$ in the limit, which is the best possible case.

The synchronization required for pipeline parallelism can be implemented using phasers by allocating an array of phasers, such that phaser `ph[i]` is “signalled” in iteration i by a call to `ph[i].arrive()` as follows:

```
1 // Code for pipeline stage i
2 while ( there is an input to be processed ) {
3     // wait for previous stage, if any
4     if ( i > 0 ) ph[i - 1].awaitAdvance();
5
6     process input;
7
8     // signal next stage
9     ph[i].arrive();
10 }
```

Optional Reading:

1. Wikipedia article on [Pipeline \(computing\)](#).